

Enron POI Classifier Model Building

Initializing and checking versions

In [1]:

```
%cd C:\Users\damia\OneDrive\Dokumente\Udacity\DataAnalyst\EnronFinal\final_project
import sys
import sklearn
print "sklearn", sklearn.__version__
print "python", sys.version
```

```
C:\Users\damia\OneDrive\Dokumente\Udacity\DataAnalyst\EnronFinal\final_pro
ject
sklearn 0.17.1
python 2.7.12 |Anaconda 4.1.1 (64-bit)| (default, Jun 29 2016, 11:07:13)
[MSC v.1500 64 bit (AMD64)]
```

Loading the cleaned data (see EDA document for details)

In [2]:

```
import sys
import pickle
import numpy as np
with open("clean_data.pkl", "r") as data_file:
    data_dict = pickle.load(data_file)
```

Collect labels and features in appropriate arrays

In [14]:

```
def init_feature_names():
    global feature_names
    feature_names = sorted(data_dict['ALLEN PHILLIP K'].keys())
    feature_names.remove('poi') # the Label
    feature_names.remove('email_address') # unique for each data point, thus exclude
    feature_names.remove('total_payments') # linear combination of other features
    feature_names.remove('total_stock_value') # linear combination of other features
def update_data():
    global labels, features
    labels = [1.0 if d['poi'] else 0.0 for d in data_dict.values()]
    features = [[float(d[f]) if f in d and d[f] != 'NaN' else 0.0 for f in feature_name
s] for d in data_dict.values()]
init_feature_names()
update_data()
```

Standardize features for use in models that are sensitive to magnitudes

In [6]:

```
def get_feature_values(fidx):  
    return [features[i][fidx] for i in range(len(features))]  
mu_f = [np.mean(get_feature_values(fidx)) for fidx in range(len(feature_names))]  
sd_f = [np.std(get_feature_values(fidx)) for fidx in range(len(feature_names))]  
norm_features = []  
for idx in range(len(features)):  
    data = []  
    for fidx in range(len(feature_names)):  
        data.append((features[idx][fidx] - mu_f[fidx]) / sd_f[fidx] if sd_f[fidx] != 0  
else features[idx][fidx])  
    norm_features.append(data)
```

Import various stuff needed for model building

In [8]:

```
from sklearn.cross_validation import StratifiedShuffleSplit  
from sklearn.grid_search import GridSearchCV
```

Decision Tree Classifier

Feature selection

In [9]:

```
from sklearn.tree import DecisionTreeClassifier
splitter = StratifiedShuffleSplit(labels, n_iter=1000, random_state=42)
scores = {}
for name in feature_names:
    scores[name] = []
for idx_train, idx_test in splitter:
    features_train = [features[i] for i in idx_train]
    features_test = [features[i] for i in idx_test]
    labels_train = [labels[i] for i in idx_train]
    labels_test = [labels[i] for i in idx_test]
    clf = DecisionTreeClassifier()
    clf.fit(features_train, labels_train)
    for i in range(len(feature_names)):
        scores[feature_names[i]].append(clf.feature_importances_[i])
for name, values in scores.items():
    scores[name] = np.mean(values)
cumulative_score = 0
for name, score in sorted(scores.items(), key=lambda x: x[1], reverse=True):
    cumulative_score += score
    print "%-25s: %.3f %.3f" % (name, score, cumulative_score)
```

```
exercised_stock_options : 0.198 0.198
fraction_to_poi          : 0.149 0.348
other                    : 0.124 0.471
shared_receipt_with_poi : 0.111 0.583
expenses                 : 0.107 0.690
bonus                    : 0.096 0.786
restricted_stock         : 0.041 0.827
long_term_incentive     : 0.036 0.863
deferred_income         : 0.036 0.899
salary                   : 0.021 0.920
from_messages           : 0.019 0.940
from_this_person_to_poi : 0.019 0.958
deferral_payments       : 0.011 0.969
fraction_from_poi        : 0.011 0.980
from_poi_to_this_person : 0.010 0.989
to_messages              : 0.009 0.998
loan_advances            : 0.001 0.999
restricted_stock_deferred: 0.001 1.000
director_fees            : 0.000 1.000
```

I select the top 5 features for the model

In [10]:

```
feature_names = ['exercised_stock_options', 'fraction_to_poi', 'other', 'shared_receipt
_with_poi', 'expenses']
update_data()
```

Algorithm Tuning

In [12]:

```
parameters = {'max_depth': (1, 2, 3, 4, 5, None),
              'min_samples_split': (2, 3, 4),
              'min_samples_leaf': (1, 2, 3, 4),
              'max_leaf_nodes': (2, 3, 4, 5, 10, None),
              'criterion': ('gini', 'entropy')}
clf = GridSearchCV(DecisionTreeClassifier(),
                  parameters,
                  scoring='f1',
                  cv=StratifiedShuffleSplit(labels, n_iter=100, random_state=42))
clf.fit(features, labels)
print 'Score: %0.3f' % clf.best_score_
print 'Parameters:'
for name, value in sorted(clf.best_estimator_.get_params().items()):
    print '%s: %r' % (name, value)
```

```
Score: 0.513
Parameters:
class_weight: None
criterion: 'entropy'
max_depth: 4
max_features: None
max_leaf_nodes: None
min_samples_leaf: 4
min_samples_split: 3
min_weight_fraction_leaf: 0.0
presort: False
random_state: None
splitter: 'best'
```

Logistic Regression (Lasso)

Feature selection

In [32]:

```
init_feature_names()
update_data()
from sklearn.linear_model import LogisticRegressionCV, LogisticRegression
clf = LogisticRegressionCV(cv=StratifiedShuffleSplit(labels, n_iter=1000,
random_state=42),
                        scoring='f1',
                        solver='liblinear',
                        penalty='l1')
clf.fit(norm_features, labels)
coefs = [(feature_names[idx], clf.coef_[0][idx]) for idx in range(len(feature_names))]
for name, value in sorted(coefs, key=lambda x: abs(x[1]), reverse=True):
    print '%-27s: %8.4f' % (name, value)
```

```
from_messages           : -31.1839
restricted_stock_deferred : 11.7534
from_this_person_to_poi : 9.6442
director_fees           : -4.8050
deferral_payments       : -3.2756
deferred_income          : -2.5946
from_poi_to_this_person  : 1.5348
salary                  : 1.1790
bonus                   : -1.1330
exercised_stock_options  : 0.8050
fraction_from_poi        : -0.8016
expenses                 : 0.7647
other                    : 0.7273
to_messages              : 0.4015
loan_advances            : -0.3558
long_term_incentive      : -0.3376
restricted_stock         : -0.1765
fraction_to_poi          : 0.0921
shared_receipt_with_poi  : 0.0394
```

We take the 9 top features as they have the largest weights

In [33]:

```
feature_names = ['from_messages', 'restricted_stock_deferred', 'from_this_person_to_poi', 'director_fees',
                 'deferral_payments', 'deferred_income', 'from_poi_to_this_person', 'salary', 'bonus']
update_data()
```

Algorithm tuning

In [35]:

```
parameters = {'solver': ('liblinear',),
              'penalty': ('l1',),
              'fit_intercept': (True, False),
              'C': (.25, .5, .9, 1., 2.),
              'tol': (1e-5, 1e-4, 1e-3)
              }
clf = GridSearchCV(LogisticRegression(),
                  parameters,
                  scoring='f1',
                  cv=StratifiedShuffleSplit(labels, n_iter=100, random_state=42))
clf.fit(norm_features, labels)
print 'Score: %0.3f' % clf.best_score_
print 'Parameters:'
for name, value in sorted(clf.best_estimator_.get_params().items()):
    print '%s: %r' % (name, value)
```

```
Score: 0.441
Parameters:
C: 0.25
class_weight: None
dual: False
fit_intercept: False
intercept_scaling: 1
max_iter: 100
multi_class: 'ovr'
n_jobs: 1
penalty: 'l1'
random_state: None
solver: 'liblinear'
tol: 1e-05
verbose: 0
warm_start: False
```

Final model selection

As the decision tree has a higher F1 score (0.513) than the logistic regression model (0.441), I choose the decision tree model as the final model. In addition to this, it uses less features.